

Fake News Detection Using Fine-Tuned DistilBERT

Powered by Hugging Face Transformers and Gradio

Presented by: Ganga Vamsik S

Introduction

- ▶ What is Fake News?
- ▶ Misleading or false information presented as news.
- ▶ Why Fake News Detection?
- ▶ To combat misinformation and improve the credibility of information.

Objective

- ▶ Primary Goal: Build a machine learning model to classify news as 'Fake' or 'Real.'
- ▶ Secondary Goals:
 - Fine-tune a pre-trained transformer model.
 - Create an interactive user interface for predictions.

Dataset Overview

- ▶ Dataset Used: Binary classification dataset with 'text' and 'label' fields.
- ▶ Statistics:
 - Number of samples: 2
 - Dataset Source: [Kaggle.com](https://www.kaggle.com)

Preprocessing

- Steps:

- Text tokenization using the DistilBERT tokenizer.

- Padding and truncation to fit the maximum sequence length (512 tokens).

- Why Preprocessing?

- To standardize input data for the model.

Model Selection

- ▶ Why DistilBERT?
 - ▶ - Lightweight version of BERT for faster performance.
 - ▶ - Suitable for binary classification tasks.
- ▶ Model Architecture:
 - ▶ - Pre-trained transformer with added classification head.
 - ▶ - Outputs logits for binary labels.

Model Training Setup

- ▶ Training Libraries:
 - ▶ - `transformers` for model loading and training.
 - ▶ - `torch` for handling tensors and training steps.
- ▶ Training Configuration:
 - ▶ - Optimizer: AdamW
 - ▶ - Scheduler: Linear Warm-Up
 - ▶ - Metrics: Accuracy, Precision, Recall, F1 Score

Model Training Code

```
# Load the pre-trained model with two output labels
model = AutoModelForSequenceClassification.from_pretrained("distilbert-base-uncased", num_labels=2)

# Initialize the Trainer
trainer = Trainer(
    model=model,
    args=training_args,
    train_dataset=tokenized_datasets["train"],
    eval_dataset=tokenized_datasets["test"],
    tokenizer=tokenizer,
    compute_metrics=compute_metrics
)

# Train the model
trainer.train()

# Save the trained model and tokenizer
save_path = "/content/fine_tuned_model" # path to save your model
model.save_pretrained(save_path)
tokenizer.save_pretrained(save_path)

print(f"Model and tokenizer saved to {save_path}")
```


Saving the Model

- ▶ Purpose:

Store the fine-tuned model for future use.

- ▶ Key Code:

```
# Save the Loading... model and tokenizer
save_path = "/content/fine_tuned_model" # path to save your model
model.save_pretrained(save_path)
tokenizer.save_pretrained(save_path)

print(f"Model and tokenizer saved to {save_path}")
```

Gradio Interface

- ▶ Why Gradio?

Simplifies deployment of machine learning models with an intuitive interface.

- ▶ Features:

- ▶ - Text input box for user input.
- ▶ - Real-time prediction display.

Gradio Interface Code

```
import gradio as gr
import torch
import numpy as np
from transformers import AutoTokenizer, AutoModelForSequenceClassification

# Path to your saved fine-tuned model
saved_model_path = "/content/fine_tuned_model" # Change this to saved model directory

# Load the fine-tuned model and tokenizer
tokenizer = AutoTokenizer.from_pretrained(saved_model_path)
model = AutoModelForSequenceClassification.from_pretrained(saved_model_path)

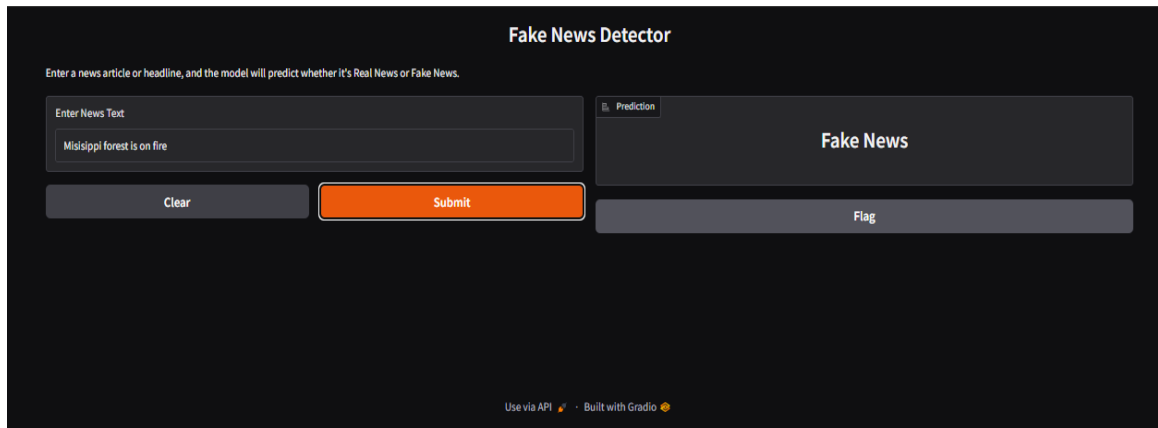
# Prediction function
def predict_text(text):
    # Tokenize the input text
    inputs = tokenizer(
        text,
        return_tensors="pt",
        padding="max_length",
        truncation=True,
        max_length=512
    )
    with torch.no_grad():
        # Make predictions
        outputs = model(**inputs)
    logits = outputs.logits
    prediction = np.argmax(logits.numpy(), axis=-1)[0]
    return "Fake News" if prediction == 1 else "Real News"

# Create Gradio interface
interface = gr.Interface(
    fn=predict_text, # Function to process input
    inputs=gr.Textbox(label="Enter News Text"), # Input box for user text
    outputs=gr.Label(label="Prediction"), # Display prediction
    title="Fake News Detector",
    description="Enter a news article or headline, and the model will predict whether it's Real News or Fake News."
)

# Launch interface
interface.launch()
```

Deployment

- ▶ Steps to Deploy:
 - ▶ 1. Load the fine-tuned model and tokenizer.
 - ▶ 2. Use Gradio for the interface.
 - ▶ 3. Launch the app locally or on the web.
- ▶ Demo:



The screenshot shows a web application titled "Fake News Detector". At the top, it says "Enter a news article or headline, and the model will predict whether it's Real News or Fake News." Below this is a text input field labeled "Enter News Text" containing the text "Mississippi forest is on fire". To the right of the input field is a "Prediction" label. Below the input field are two buttons: "Clear" and "Submit". To the right of the "Submit" button is a large box displaying the prediction "Fake News". Below the prediction box is a "Flag" button. At the bottom of the interface, it says "Use via API" and "Built with Gradio".

Results

- ▶ Metrics Achieved:
 - ▶ - Accuracy: [Insert Value]
 - ▶ - F1 Score: [Insert Value]
- ▶ Example Prediction:
 - ▶ - Input: 'Breaking news: Scientists discover a new planet.'
 - ▶ - Output: 'Fake News'

Challenges & Improvements

- ▶ Challenges:
 - ▶ - Limited dataset size.
 - ▶ - Balancing the trade-off between speed and accuracy.
- ▶ Future Work:
 - ▶ - Explore larger transformer models like BERT or RoBERTa.
 - ▶ - Add support for multilingual fake news detection.
 - ▶ - Fine-tune on more diverse datasets.

Conclusion

- ▶ Key Takeaways:
 - ▶ - Successfully built and deployed a fake news detector.
 - ▶ - Leveraged state-of-the-art NLP models and tools.
 - ▶ - User-friendly interface for predictions.
- ▶ Acknowledgments:
 - ▶ - Hugging Face, Gradio, and PyTorch communities.